

Project 3 Report

Yizhan Ao

Email: josephao@umd.edu
UID: 116022064

Yingqiao Gou

Email: ygou@terpmail.umd.edu
UID: 115979000

Abstract

You will be graded primarily based on your report. The most important part is to understand and explain why your algorithm is working or not working on a given dataset. What do you think are the drawbacks of your code!?

We want you to demonstrate an understanding of the concepts involved in the project, and to show the output produced by your code. Include visualizations of the output of each stage in your pipeline (as shown in the system diagram on page 2), and a description of what you did for each step. Assume that we're familiar with the project, so you don't need to spend time repeating what's already in the course notes. Instead, focus on any interesting problems you encountered and/or solutions you implemented.

1 Concept

The aim of this project is to segment deformable object from a given video sequence. This document just provides an overview of what you need to do. For a full breakdown of how each step in the pipeline works, see the course notes for this project.

2 Setting up local windows

Finding and saving the window's midpoints was all that was required to initialize the local windows. These mid-points were just points uniformly dispersed over the first mask's shape. InitLocalWindows.m has previously done this for us. We didn't alter anything except the problem that was in the code at the time.

3 Initializes color models

We used rgb2lab to convert the original frame from RGB to LAB colorspace in order to initialize the color models. Then, using bwdist, to calculate the distance matrix of the mask outline in each window. This provides me the distance to the original mask's edge, which allows us to eliminate points that are too near to the edge from GMM training since color changes dramatically at this edge and may not be representational of either the foreground or background.

$$p_c(x) = p_c(x \mid \mathcal{F}) / (p_c(x \mid \mathcal{F}) + p_c(x \mid \mathcal{B})) \quad (1)$$

The color model is started by estimating a Gaussian Mixture Model for each window's foreground pixels, and then computing a GMM for the background pixels (where the mask is produced by the user's roipoly result). Because there is little change in the colors of the foreground and background in the instance of the turtle, we utilized GMM of size 1; the number k varied for the other sets of photos.

4 Initializes color confidences

The color confidence is derived using the Rotobrush paper's algorithm and is dependent on the color model. Every window's color confidences are saved. [0,1] is a nice method to represent how effectively the foreground and background are separated.

$$f_c = 1 - \frac{\int_{W_k} |L^t(x) - p_c(x)| \cdot \omega_c(x) dx}{\int_{W_k} \omega_c(x) dx} \quad (2)$$

5 Shapes Models

$$f_s(x) = 1 - \exp(-d^2(x)/\sigma_s^2) \quad (3)$$

To update foreground GMM, we sampled all pixels in the warped window whose foreground confidence estimated from the update shape model is greater than a high threshold 0.75, according to the paper. Because ShapeConfidence simply keeps track of the confident score in relation to the windows, When the window is deformed, it is updated. To train a GMM model for the current frame, we only need to cycle through every pixel in every previous warped window that has a shape confidence score greater than 0.75 to feed into the GMM model.

5.1 How to Update Shape Model

$$\sigma_s = \begin{cases} \sigma_{\min} + a(f_c - f_{\text{cutoff}})^r & f_{\text{cutoff}} < f_c \leq 1 \\ \sigma_{\min} & 0 \leq f_c \leq f_{\text{cutoff}} \end{cases} \quad (4)$$

The paper discusses the two sigma cases in Formula 4. I assume that has been addressed in the function because the shape model is dependent on the `initShapeConfidences()` Then we do simply compare the total number of foreground pixels in the prior and current GMMs after the update.

```
for (every shape confidences):
    use color confidences and sigma to do some calculations...
    ShapeConfidences.Confidences{i} = new_f_s;
    % This is from the Piazza that is very helpful to us!
```

5.2 How to find the foreground probability of each window

$$p_{\mathcal{F}}^k(x) = f_s(x)L^{t+1}(x) + (1 - f_s(x))p_c(x) \quad (5)$$

Every window must be iterated through to determine the foreground probability of each pixel. The $pkF(x)$ cellular array would be a `1xWindows` count cellular array that stores the pixels' `WidthxWidth` foreground probability. When we utilize the warped mask here, we must first remove the appropriate patch from the current window.

```
current_mask_window = warpedMask(current_window_xRange , current_window_yRange);
```

6 Updating Window Locations

We utilized `detectSURFFeatures`, which is rotational invariant, to estimate the vast quantities of motion in the object. We tried deleting the backdrop (setting pixels to NaN) to compel matching to focus on the foreground, however this frequently resulted in the algorithm failing to discover enough matching sites. Even with an iterative matching threshold, we were unable to resolve this issue. Even if there were matches early on, the algorithm would eventually fail. This, we believe, is the major source of the algorithm's failure in rotating frame sets. Some of the matches are fine, while some are terrible, resulting in a bad transformation, as illustrated below. We think one of the things we shouldn't spend too much time on is that twerking the parameters(including the size of the window and numbers of the windows). I think we have tried too hard to find the best result of the window positioning.

7 Calculates local window movement based on optical flow between frames

This step is to estimate the local boundry and deform the picture. The transformation was insufficient to track modest amounts of motion. To account for these little variations, we employed optical flow. We determined the average optical flow in the X and Y directions for the foreground pixels in each window and added it to the points in that window.

8 Finds affine transform between two frames.

9 Update shape and color models.

10 Work Distribution

10.1 Yizhan Ao

Setup Local Windows + Initialize Color Models + Combine Shape and Color Models + Merge Local Windows + Extract final foreground mask + Update Color Model (and color confidence)+ Report(1-6)

10.2 Yingqiao Gou

Compute Color Model Confidence + Initialize Shape Model + Compute Shape confidence + Estimate Entire-Object Motion + Estimate Local Boundary Deformation(6-12)

11 Output

12 Conclusion

You have been provided the frames from five video clips. Run your code on each set of frames and create videos, named as indicated above. For each video track the following:

Frames1: Track the turtle. Source: An underwater video captured by Chahat at Lakshadweep islands, India.

Frames2: Track the motorcycle and rider. Source.

Frames3: Track the gymnast. Source.

Frames4: Track the powerlifter, without the weights. Source.

Frames5: Track the lizard (including tail). Source.

References

- [1] Bai, X., Wang, J., Simons, D. and Sapiro, G., 2009, July. Video snapcut: robust video object cutout using localized classifiers. In ACM Transactions on Graphics (ToG) (Vol. 28, No. 3, p. 70). ACM.
- [2] Wang, J. and Cohen, M.F., 2005, October. An iterative optimization approach for unified image segmentation and matting. In Computer Vision, 2005. ICCV 2005. Tenth IEEE International Conference on (Vol. 2, pp. 936-943). IEEE.